

Assignment-4

Assignment Date	17 October 2022
Student Name	V.Kalpana
Student Roll Number	962719106012
Maximum Marks	2Marks

Question 1:

Download the dataset

Link:

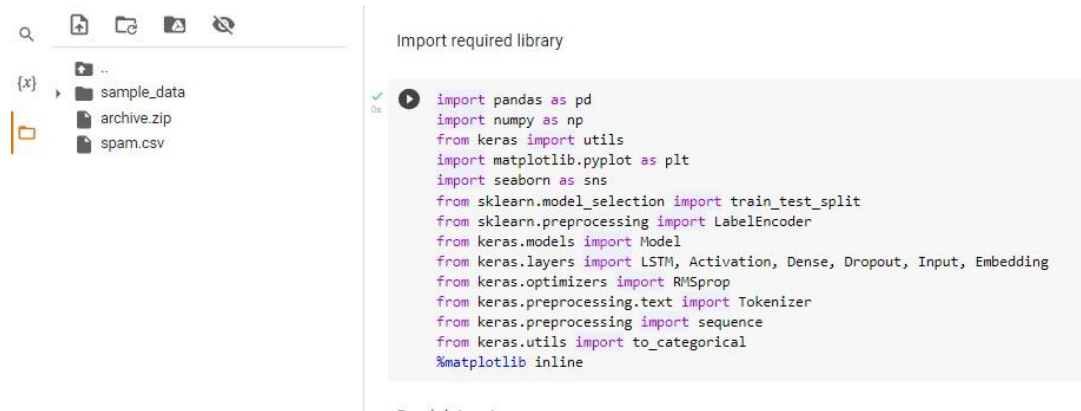
https://drive.google.com/file/d/1Sjqx5H5R86tRp2YZKzzd4_iEfjChZ3ob/view?usp=sharing

Question 2:

Import required library

Solution:

```
import pandas as pd
import numpy as np
from keras import utils
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import to_categorical
```



Question 3:

Read dataset and do pre-processing

Solution:

Read dataset

```
!unzip "/content/archive.zip"
```

```
df = pd.read_csv('spam.csv', delimiter=',', encoding='latin-1')
df
```

Pre processing

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)
df
```

```
sns.countplot(df.v1, palette='Set3')
plt.xlabel('Label')
plt.title('Number of ham and spam messages')
```

```
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.15)
```

```
max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = utils.pad_sequences(sequences, maxlen=max_len)
```

```
sequences_matrix = np.reshape(sequences_matrix,(4736,150,1))
sequences_matrix.ndim
```

The screenshot shows a Jupyter Notebook environment. On the left is a file explorer with a folder named 'sample_data' containing 'archive.zip' and 'spam.csv'. The main area displays two code cells. The first cell runs `!unzip "/content/archive.zip"`, showing output for unzipping 'spam.csv'. The second cell runs `df = pd.read_csv('spam.csv', delimiter=',', encoding='latin-1')`. Below the code is a preview of the 'df' DataFrame, showing columns 'v1' and 'v2' with text data, and three unnamed columns containing 'NaN' values.

The screenshot shows a Jupyter Notebook interface. On the left, a file explorer displays a directory containing 'sample_data', 'archive.zip', and 'spam.csv'. The main area shows a code cell where the command `df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)` has been executed successfully, indicated by a green checkmark and '2s' runtime. Below the code, a preview of the resulting DataFrame is shown as a table:

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will i_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

At the bottom, it states '5572 rows x 2 columns'.



```
[7] X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)

[8] X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)

[9] max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = utils.pad_sequences(sequences,maxlen=max_len)

[10] sequences_matrix.shape
(4736, 150)

[11] sequences_matrix.ndim
2

[12] sequences_matrix = np.reshape(sequences_matrix,(4736,150,1))
sequences_matrix.ndim
3
```

Question 4:

Create model

Solution:

```
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers import Embedding
```

```
model = Sequential()
```

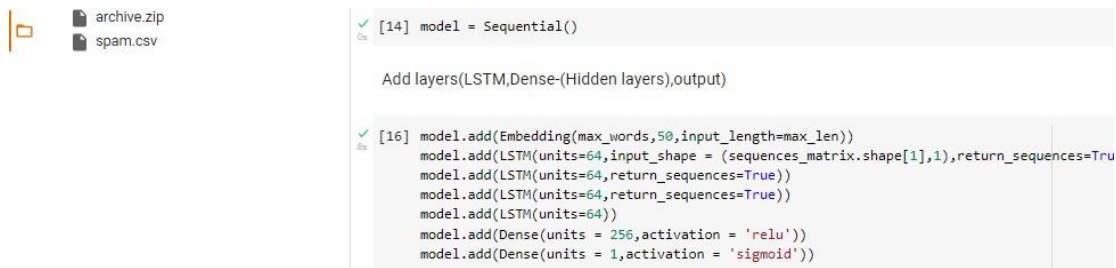


Question 5:

Add layers(LSTM,Dense-(Hidden layers),output)

Solution:

```
model.add(Embedding(max_words,50,input_length=max_len))
model.add(LSTM(units=64,input_shape = (sequences_matrix.shape[1],1),return_sequences=True))
model.add(LSTM(units=64,return_sequences=True))
model.add(LSTM(units=64,return_sequences=True))
model.add(LSTM(units=64))
model.add(Dense(units = 256,activation = 'relu'))
model.add(Dense(units = 1,activation = 'sigmoid'))
```



Question 6:

Compile the model

Solution:

```
model.summary()
model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])
```

{x}
sample_data
archive.zip
spam.csv

Compile the model

```
[17] model.summary()
model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 150, 64)	29440
lstm_1 (LSTM)	(None, 150, 64)	33024
lstm_2 (LSTM)	(None, 150, 64)	33024
lstm_3 (LSTM)	(None, 64)	33024
dense (Dense)	(None, 256)	16640
dense_1 (Dense)	(None, 1)	257

```

Total params: 195,409
Trainable params: 195,409
Non-trainable params: 0

```

Question 7:

Fit the model

Solution:

X = model.fit(sequences_matrix,Y_train,batch_size=128,epochs=5,validation_split=0.2)
X

{x}
sample_data
archive.zip
spam.csv

Fit the model

```
[18] X = model.fit(sequences_matrix,Y_train,batch_size=128,epochs=5,validation_split=0.2)
X
```

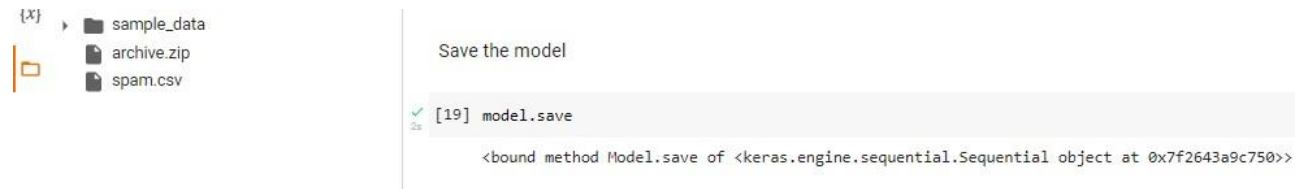
Epoch 1/5
30/30 [=====] - 43s 1s/step - loss: 0.4490 - accuracy: 0.8688 - val_loss: 0.4257 - val_accuracy: 0.8513
Epoch 2/5
30/30 [=====] - 33s 1s/step - loss: 0.2615 - accuracy: 0.9092 - val_loss: 0.1283 - val_accuracy: 0.9610
Epoch 3/5
30/30 [=====] - 35s 1s/step - loss: 0.0724 - accuracy: 0.9794 - val_loss: 0.0096 - val_accuracy: 0.9705
Epoch 4/5
30/30 [=====] - 33s 1s/step - loss: 0.0529 - accuracy: 0.9857 - val_loss: 0.0053 - val_accuracy: 0.9778
Epoch 5/5
30/30 [=====] - 33s 1s/step - loss: 0.0392 - accuracy: 0.9900 - val_loss: 0.0036 - val_accuracy: 0.9768
<keras.callbacks.History at 0x7f263f739bd0>

Question 8:

Save the model

Solution:

model.save



Question 9:

Test the model

Solution:

```
test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix = utils.pad_sequences(test_sequences,maxlen=max_len)
```

```
accr = model.evaluate(test_sequences_matrix,Y_test)
```

```
l = accr[0]
a =accr[1]
print('Test set\n Loss: {:.3f}\n Accuracy: {:.3f}'.format(l,a))
```

